

Multiparty Computation made Practical

Using the Virtual Ideal Functionality Framework

Martin Geisler
⟨mg@daimi.au.dk⟩

BRICS
Department of Computer Science
University of Aarhus
Denmark

June 19, 2008

Outline

Background

- Multiparty Computation

- Related Projects

Design Goals

- Automatic Parallel Execution

- Program Counters

Benchmark Results

- Multiplications

- Comparisons

Possible Improvements


- Program Counters

- Memory Overheads

- Debugging


Conclusion

Multiparty Computation




- ▶ n players
- ▶ Wish to jointly compute f
- ▶ Player P_i has input x_i
- ▶ Players learn
 $y = f(x_1, x_2, \dots, x_n)$

Multiparty Computation




- ▶ n players
- ▶ Wish to jointly compute f
- ▶ Player P_i has input x_i
- ▶ Players learn
 $y = f(x_1, x_2, \dots, x_n)$
- ▶ Up to t players are **corrupt**
- ▶ Must keep inputs **private**
- ▶ Players **only** learn y

Trivial “Solution”




- Send inputs to trusted party over secure channels

Trivial “Solution”




- ▶ Send inputs to trusted party over secure channels
- ▶ T computes $y = f(x_1, x_2, \dots, x_n)$

Trivial “Solution”



- ▶ Send inputs to trusted party over secure channels
- ▶ T computes
$$y = f(x_1, x_2, \dots, x_n)$$
- ▶ Distributes y to all players


Trivial “Solution”




- ▶ Send inputs to trusted party over secure channels
- ▶ T computes
$$y = f(x_1, x_2, \dots, x_n)$$
- ▶ Distributes y to all players
- ▶ Obviously secure!
- ▶ But **who** should play T ?

Security Model: UC Framework

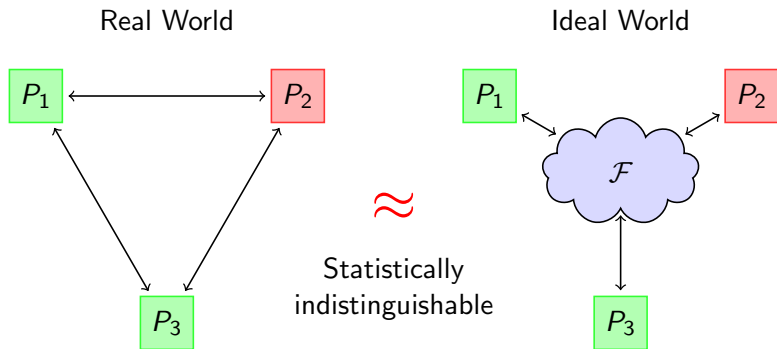
Real World



Ideal World



Security Model: UC Framework




Functionality provided by VIFF

Supported operations:

- ▶ Input
- ▶ Output
- ▶ Addition
- ▶ Multiplication
- ▶ Comparison

All over arbitrary finite fields.



Related Projects

- ▶ FairPlay project (Haifa, Israel)
 - ▶ Yao-garbled circuits for 2 players
 - ▶ Java implementation
 - ▶ Working on multiparty computation

Related Projects

- ▶ FairPlay project (Haifa, Israel)
 - ▶ Yao-garbled circuits for 2 players
 - ▶ Java implementation
 - ▶ Working on multiparty computation
- ▶ SCET project (Aarhus, Denmark)
 - ▶ General multiparty computations
 - ▶ C# implementation

Related Projects

- ▶ FairPlay project (Haifa, Israel)
 - ▶ Yao-garbled circuits for 2 players
 - ▶ Java implementation
 - ▶ Working on multiparty computation
- ▶ SCET project (Aarhus, Denmark)
 - ▶ General multiparty computations
 - ▶ C# implementation
- ▶ SIMAP project (Aarhus, Denmark)
 - ▶ General multiparty computations
 - ▶ Java implementation
 - ▶ Some work done on a domain specific language

Outline

Background

Multiparty Computation

Related Projects

Design Goals

Automatic Parallel Execution

Program Counters

Benchmark Results

Multiplications

Comparisons

Possible Improvements

Program Counters

Memory Overheads

Debugging

Conclusion

VIFF Goals

- ▶ Easy to use for the programmer:

```
x = 2 * a - b * c
```


VIFF Goals

- ▶ Easy to use for the programmer:

```
x = 2 * a - b * c
```

- ▶ Automatically run things in parallel:

```
x = a * b  
y = c * d  
z = e * f
```

VIFF Goals

- ▶ Easy to use for the programmer:

```
x = 2 * a - b * c
```

- ▶ Automatically run things in parallel:

```
x = a * b  
y = c * d  
z = e * f
```

- ▶ Extensible with new operations:

```
def max(a, b):  
    c = a > b  
    return c * a + (1 - c) * b
```

Parallel Execution

- ▶ Networks have significant latency
- ▶ Want to run many operations in parallel
- ▶ Including primitive and compound operations

Example: Hamming Distance

```
def xor(a, b):  
    assert a.field is b.field  
    if a.field is GF256:  
        return a + b  
    else:  
        return a + b - 2 * a * b
```

- ▶ Straight-forward exclusive-or
- ▶ Fast for $GF(2^8)$ elements
- ▶ Slower for \mathbb{Z}_p elements
- ▶ (Already part of VIFF)

Example: Hamming Distance


```
def xor(a, b):  
    assert a.field is b.field  
    if a.field is GF256:  
        return a + b  
    else:  
        return a + b - 2 * a * b
```

```
def hamming(s, t):  
    distance = 0  
    for i in range(len(s)):  
        distance += xor(s[i], t[i])  
    return distance
```

- ▶ Straight-forward exclusive-or
 - ▶ Fast for $GF(2^8)$ elements
 - ▶ Slower for \mathbb{Z}_p elements
 - ▶ (Already part of VIFF)
-
- ▶ Hamming distance
 - ▶ xor calls should run in parallel


Hamming Distance Execution Tree

hamming(s, t) translates to:




Hamming Distance Execution Tree

hamming(s, t) translates to:




xor(a, b) translates to:



VIFF Execution Strategy

- ▶ Create execution tree as we go along
- ▶ Destroy execution tree from bottom up
 - ▶ Each node waits on nodes below
 - ▶ Bottom nodes wait on network traffic
- ▶ **Composable**: just plug new operations into tree!

Program Counters



- ▶ No parsing — execution tree never fully constructed
- ▶ No fixed evaluation order
- ▶ But we must identify results

Program Counters Implementation

- ▶ First attempt:
 - ▶ Each player increments a global counter
 - ▶ Fails because of asynchronous execution

Program Counters Implementation

- ▶ First attempt:
 - ▶ Each player increments a global counter
 - ▶ Fails because of asynchronous execution
- ▶ Working solution:
 - ▶ Manually “weave” a **program counter** through program
 - ▶ Tedious, easy to forget to increment program counter

Program Counters Implementation

- ▶ First attempt:
 - ▶ Each player increments a global counter
 - ▶ Fails because of asynchronous execution
- ▶ Working solution:
 - ▶ Manually “weave” a **program counter** through program
 - ▶ Tedious, easy to forget to increment program counter
- ▶ Current solution:
 - ▶ Runtime methods wrapped by a decorator
 - ▶ Calculated automatically based on call stack

Outline

Background

- Multiparty Computation

- Related Projects

Design Goals

- Automatic Parallel Execution

- Program Counters

Benchmark Results

- Multiplications

- Comparisons

Possible Improvements

- Program Counters

- Memory Overheads

- Debugging

Conclusion

Performance Results

- ▶ Tested on 3 machines: USA, Norway, and Denmark
- ▶ Results for VIFF 0.4 (VIFF 0.6 is similar or better)
- ▶ Tested multiplications and comparisons

Performance Results

- ▶ Tested on 3 machines: USA, Norway, and Denmark
- ▶ Results for VIFF 0.4 (VIFF 0.6 is similar or better)
- ▶ Tested multiplications and comparisons
- ▶ Tested parallel and serial multiplications

A



B



A




B





Performance Results

- ▶ Tested on 3 machines: USA, Norway, and Denmark
- ▶ Results for VIFF 0.4 (VIFF 0.6 is similar or better)
- ▶ Tested multiplications and comparisons
- ▶ Tested parallel and serial multiplications





Performance Results

- ▶ Tested on 3 machines: USA, Norway, and Denmark
- ▶ Results for VIFF 0.4 (VIFF 0.6 is similar or better)
- ▶ Tested multiplications and comparisons
- ▶ Tested parallel and serial multiplications





Performance Results

- ▶ Tested on 3 machines: USA, Norway, and Denmark
- ▶ Results for VIFF 0.4 (VIFF 0.6 is similar or better)
- ▶ Tested multiplications and comparisons
- ▶ Tested parallel and serial multiplications




Performance Results

- ▶ Tested on 3 machines: USA, Norway, and Denmark
- ▶ Results for VIFF 0.4 (VIFF 0.6 is similar or better)
- ▶ Tested multiplications and comparisons
- ▶ Tested parallel and serial multiplications




Parallel Multiplications

Multiplying random 65-bit numbers:




Serial Multiplications

Multiplying random 65-bit numbers:



Parallel Comparisons

Comparing random 32-bit numbers, 65-bit field modulus:



Outline

Background

- Multiparty Computation

- Related Projects

Design Goals

- Automatic Parallel Execution

- Program Counters

Benchmark Results

- Multiplications

- Comparisons

Possible Improvements

- Program Counters

- Memory Overheads

- Debugging

Conclusion

Program Counter Overheads

- ▶ They work, but it's a bit magic...
- ▶ Exactly when must the program counter be updated?
- ▶ Excessive wrapping slows down method calls
- ▶ Program counter size depends on stack depth

Memory Overheads

- ▶ Python objects have a large memory overhead:
 - ▶ Reference count (4 bytes)
 - ▶ Object attribute dictionary (144 bytes)
 - ▶ Allocation overhead (>8 bytes)
 - ▶ ...

Memory Overheads

- ▶ Python objects have a large memory overhead:
 - ▶ Reference count (4 bytes)
 - ▶ Object attribute dictionary (144 bytes)
 - ▶ Allocation overhead (>8 bytes)
 - ▶ ...
- ▶ 100,000 field elements with 65-bit prime:
 - ▶ Optimal: ≈ 800 KB
 - ▶ VIFF: ≈ 40 MB (expanded 50 times)
 - ▶ More memory needed for execution tree

Debugging

- ▶ Something went wrong! What now?
- ▶ Debugging asynchronous programs can be hard

Debugging

- ▶ Something went wrong! What now?
- ▶ Debugging asynchronous programs can be hard
- ▶ Need better
 - ▶ Logging infrastructure
 - ▶ Handling of exceptions

Type Safety

- ▶ Python is a strongly typed language
 - ▶ `"12" * 3 == "121212"` but `12 * 3 == 36`
 - ▶ `"1" + 12` raises `TypeError`
- ▶ Types are only checked at runtime

Type Safety

- ▶ Python is a strongly typed language
 - ▶ `"12" * 3 == "121212"` but `12 * 3 == 36`
 - ▶ `"1" + 12` raises `TypeError`
- ▶ Types are only checked at runtime
- ▶ Unit tests help here
- ▶ Better input validation

Outline

Background

- Multiparty Computation

- Related Projects

Design Goals

- Automatic Parallel Execution

- Program Counters

Benchmark Results

- Multiplications

- Comparisons

Possible Improvements

- Program Counters

- Memory Overheads

- Debugging

Conclusion

Conclusion

- ▶ VIFF provides a general framework for MPC
- ▶ Automatic parallel execution
- ▶ Free Software: LGPL
- ▶ Please see <http://viff.dk/>
 - ▶ Source code
 - ▶ Documentation
 - ▶ Mailing list

Conclusion

- ▶ VIFF provides a general framework for MPC
- ▶ Automatic parallel execution
- ▶ Free Software: LGPL
- ▶ Please see <http://viff.dk/>
 - ▶ Source code
 - ▶ Documentation
 - ▶ Mailing list

Thank you for listening!